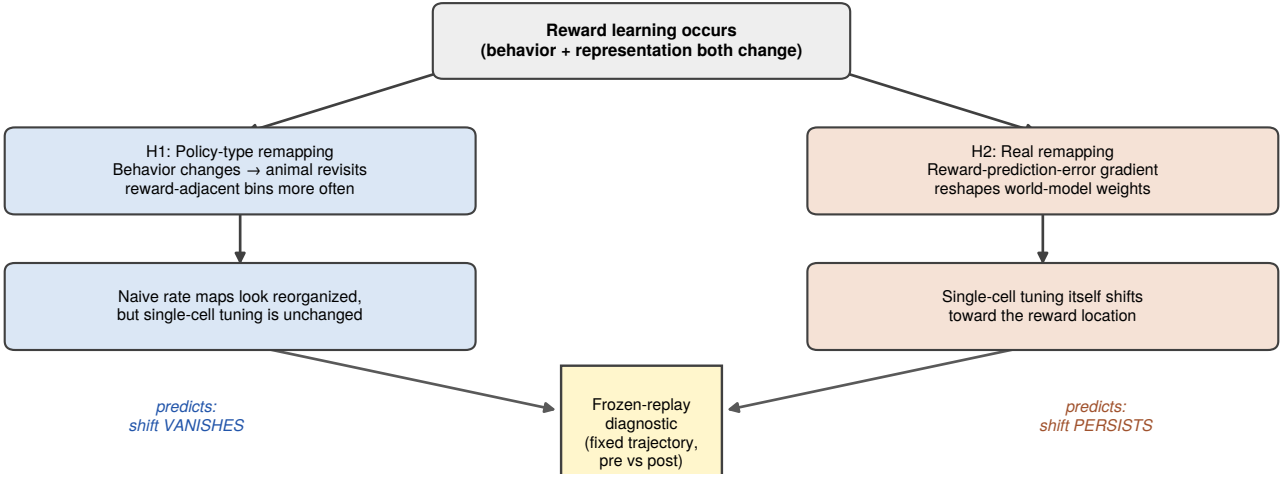
Reward learning occurs
(behavior + representation both change)
H1: Policy-type remapping
Behavior changes → animal revisits
reward-adjacent bins more often
H2: Real remapping
Reward-prediction-error gradient
reshapes world-model weights
Naive rate maps look reorganized,
but single-cell tuning is unchanged
Single-cell tuning itself shifts
toward the reward location
predicts:
shift VANISHES
Frozen-replay
diagnostic
(fixed trajectory,
pre vs post)
predicts:
shift PERSISTS